

Web Data Management -Programming Project

Group 8

Georgios Dimitropoulos, 4727657
Manolis Manousogiannis, 4727517
Georgia Zarnomitrou, 4724089

1 INTRODUCTION

In this paper we present the inner workings of four scalable database systems (PostgreSQL, Apache Spark, Apache Lucene and Neo4j) with the help of a movie database created with information gathered from IMDB. Specifically, we display the strengths and weaknesses of each system and the lessons we have learned from our attempts to implement various queries in each system.

2 DATASET

In order to examine the performance of the four aforementioned systems we worked with a database regarding movies and tv series. The information in this database was gathered from IMDB and given to us by professor Lofi in the form of a backup file from a postgresQL database. This database contains a variety of information like the titles of the movies or tv series, the names of the actors and the characters they portrayed, plot keywords, the name of the producing companies and other miscellaneous data. This database is composed of 18 tables where each table has hundreds of thousands or millions of rows and has a size of 6.3GB. The size of our database proved to be a problem during our implementations since we often needed to convert it online in other formats for use in different systems and some systems could not handle such large files.

Firstly, we attempted to merge this database with another one we found online[1] because this smaller dataset contained additional information regarding the success of some movies and the popularity of its actors and cast. However this proved an impossible task due to PostgreSQL's lack of sufficient built-in functions to handle strings, so we decided to add this additional information in the form of fictitious data. For this purpose we wrote our own scripts in PostgreSQL (these scripts along with our code for each system can be found here [2]) that created columns for the budget, gross and duration of the movies in the table movie_info. We also added a column in the table cast_info to present the likes the cast of a movie has in Facebook and one column in the table aka_name that represents the likes each actor has in Facebook.

3 QUERIES

In this section we present the various queries we want to implement in the four database systems we have chosen.

- (1) Find the ratio of the facebook likes of a specific actor to those of the cast in all the movies he has participated in.
- (2) Movie titles where the director is Abd Al-Hamid, Ja 'far and the actor is "Tom Hanks".
- (3) Return the number of companies who have produced films with the word love in the title.

- (4) The titles of the 1000 more successful movies from 1990 until 2000.
- (5) The titles of the movies where an actor and the character he/she portrays have the same name.
- (6) Find the title of each movie that contains the word "design" in the list of plot keywords.
- (7) How many times does a word appear in the whole database.
- (8) Return the sum of movie gross of the movies produced by "Amazon" for different production years.
- (9) Return average of the movie budget spent by the company "Paramount Television" for different production years.
- (10) How many times does the name "Harry" appear in all the TV series of the database.
- (11) Find relation between tv series success and number of episodes.
- (12) How could we be constantly aware of the number of rows of the table titles.
- (13) Return the count of times the keyword "kill" appears in the titles of our movies.
- (14) Return the ids of keyword table tuples where keyword is equal to "kill" or phonetic code is equal to "A2525".
- (15) Return number of tuples in table keywords where keyword field starts with the word kill.

4 SYSTEM POSTGRESQL

In this section we present the system PostgreSQL and how it performed when we run the queries described in section 3. PostgreSQL is a powerful open source object-relational database that can be used to perform complex queries. The choice to pick this as one of our systems was an easy one since PostgreSQL is one of the most robust, reliable and widely known systems for handling relational databases.

In table 1 we display the execution time and planning time PostgreSQL required to perform the aforementioned queries. Planning time refers to PostgreSQL's functionality **query plan** where the system receives a query and searches for the optimal way to execute the user's request (this can be seen in figures 1 and 2 as PostgreSQL offers the possibility to view this information). In this table we can observe three empty rows, these represent queries we could not implement in PostgreSQL. Query 2 could not be implemented because of the way our database is structured, specifically the complexity of the database impedes the creation of a simple query that can go through all the necessary tables to answer the question. This shows a weakness of the Postgres system since queries depend on the way the database is structured, which is often flawed due to the human factor.

In order to resolve query 7 we would have to go through all the columns of all the tables of our database. This is an inefficient,




Figure 4: Average of movie budget spent by Paramount Television per Production Year




Figure 5: Number of times a character name (i.e. Harry) appears in all TV series

In Table 2 we can observe the execution time of our queries for the Apache Spark. As mentioned before, empty rows represent queries we could not implement. Query 2 was not implemented due to the structure of the database. Queries 5 and 11 were not able to be executed due to the size of our database and the limited resources of our computers (specifically we got an out of memory warning). Also, in order to execute query 7 we would have to go through all the columns of all the tables of our database and we found again that this was an inefficient, difficult and consuming procedure that requires again a lot of memory. Finally, we were not able to find an efficient way in order to implement query 12 that uses a trigger. However, if we were able to have more resources we could get a response for these queries.

6 SYSTEM APACHE LUCENE

In this section, we are going to present our third system, which is Apache Lucene. Lucene is an *inverted index* based text search engine library which is based entirely in Java. However, implementations

Queries	Execution Time (seconds)
1	185
2	Could not implement
3	22.5
4	270
5	Could not implement
6	41.6
7	Could not implement
8	237
9	244
10	473
11	Could not implement
12	Could not implement
13	1.42
14	3.02
15	1.49

Table 2: Execution Time of the queries for the Apache Spark

in other languages are also available and we decided to implement our system in .NET.

The reason for selecting Lucene, is that it is considered to be a really fast search engine for text due to adding content to a full-text index, so we wanted to compare its performance in this task with PostgreSQL and Spark.

In order to explain the 'inverted index' term, we could say that it is equivalent to retrieving pages of a book with a specific keyword, by searching the index at the end of the book, rather than searching each page separately.

While implementing Lucene in .NET, we went through the whole procedure that takes place in this system. As the 'unit' of index and search in Lucene is a document, we had to process our dataset as a set of Documents which contain different fields. Every document was a row of our table and every feature (row) was a field in that document. A general overview of the process and indexing of our data is shown below.




Figure 6: Process and index of data in Lucene

As Lucene is used for text search it was reasonable to test it using tables that contained text in it. From them, the most suitable was a table containing 236,000 different movie ids with their corresponding set of keywords and their phonetic code. So this means that we created 236,000 documents with 3 fields each.

After our indexing was complete we had to create a QueryParser object, which we used to pass our query to the Index searcher and retrieve our results.

we were not able to execute the queries of all systems in the same machine and this is something that justifies the slower performance of Apache Spark in some of the other queries.

In regards to the PostgreSQL system we observed that it was the only system able to handle a diverse type of queries with good performance (especially after indexing the database, the execution time was vastly improved). However it can only be used with bounded relational databases and this does not correspond well on the real world where data is generated every minute. Therefore we have reached the conclusion that all of these systems can become the ideal choice depending on the user's needs.

9 REFLECTIONS

In this section we will present the reflection we had regarding each of the systems as well as the project in general.

We have to mention that we spent a significant amount of time in trying to work with *Apache Flink*. We wanted to use this stream processing system, in order to compare it with Spark which performs better as a batch processing system. Of course the kind of data that we had did not help us in this direction as we had to split it into streams. Apart from that however, we generally believe that Flink is not easy to use if someone has no previous experience in this system. In order to use it, we had to create executable .Jar files using Java which was not easy at all. Even when more experienced with Java members of our team tried to figure it out, there was no result. We found out that there is an API() available in Python too, but the documentation and the material available online was very limited.

Regarding Spark, we did most of the work using Python, instead of Java which is mostly used, as we were more comfortable with it. However, we had many difficulties in the beginning in implementing most queries in Sql in a Python language environment.

Another major issue we had was related to our dataset. In *Lucene* for instance, that we implemented in .Net, we had to convert our .csv tables in .xml format in order to load them! Converting them was not a very big issue, but their size increased significantly, which made it impossible to use in our local machines. This could be resolved *if we used cloud services* provided from Amazon, however we did not predict this need. Using cloud services to test our platforms, would even make the comparison between the execution times of our systems more reliable as we would be able to run all systems and queries with the same computing power, rather than using different computers.

Finally, we should say that we really liked Neo4J graph processing system as it can create really useful patterns in a user friendly environment. However, we did not have the time to further discover its capabilities and create something more than those quite simple relationships we presented in the previous section.

REFERENCES

- [1] <https://data.world/data-society/imdb-5000-movie-dataset>
- [2] https://github.com/GZ-02/Web_Data_Management_Project